

The curiosity trick for solving explore-exploit problems.

Erik J Peterson^{a,b,1} and Timothy D Verstynen^{a,b,c,d}

^aDepartment of Psychology; ^bCenter for the Neural Basis of Cognition; ^cCarnegie Mellon Neuroscience Institute; ^dBiomedical Engineering, Carnegie Mellon University, Pittsburgh PA

Explore-exploit problems are common but intractable problems in reward learning, and therefore in the biological sciences. In this paper, we provide an alternative view of these problems that does not depend on exploring to optimize for reward value, even though this remains the goal of exploitation. Here exploration has no objective but learning itself, which we equate to curiosity. Through theory and experiments we show this “curiosity trick” leads to optimal value solutions for potentially all explore-exploit problems, if that is curiosity is just as important as reward collection for survival. In particular, we show our approach succeeds in naturalistic conditions when rewards are sparse, deceptive, and non-stationary. We also suggest three experiments that might be used to distinguish our approach from other theories.

Introduction

Exploration behavior can have two very different explanations depending on whether an animal expects a reward. If there is no reason to expect a reward, exploration is treated as a search for information, or as curiosity (? ? ? ? ? ?). For example, when a rat is placed in a novel environment it will explore even if no food and water is expected (? ?). If reward is expected exploration is instead interpreted as a search for reward (? ? ? ? ? ?).

An open problem in the broad field of the decision sciences is to unify exploration with exploitation, which we define as the policy of choosing the most rewarding action. When exploration optimizes for reward value, this union leads to the famous exploration-exploitation dilemma (? ? ? ? ? ?). We illustrate this classic problem in Fig. 1a.

In this paper, we offer an alternative approach. We unify exploitation with curiosity. We show that exploration is well handled by curiosity, even when the goal is to collect the most reward. We illustrate our alternative in Fig. 1b.

We can justify our use of curiosity in this alternative, because curiosity is a primary drive in most, if not all, animals (? ? ?). It is a drive that is as strong, if not sometimes stronger, than the drive for reward (? ? ? ? ? ?).

The first half of this paper we develop a the theory behind our union. This consists of mathematical study of curiosity alone, and our mathematical solution to unifying curiosity with reward collection. The second half consists of some initial simulation where we demonstrate performance which is as good, if not better, than more standard approaches. We conclude with some questions and answers about our approach.

Results

Theory. For our theoretical work, we first introduce reinforcement learning, and its basis in the Bellman equation. Next




Fig. 1. Two views of exploration and exploitation. **a.** The dilemma: either exploit an action with a known reward (e.g., return to the previous flower) or explore other actions on the chance they will return a better outcome. The central challenge here is that the outcome of exploration is uncertain, and filled with questions. **b.** An alternative view of the dilemma, with two goals: either maximize rewards or maximize information value, with a curious search of the environment. *Artist credit:* Richard Grant.

we develop a new mathematical view of information value, curiosity, and a new understanding of ideal curious search. We then use these results to show exploration-for-curiosity leads to optimal value solutions for all explore-exploit problems, with deterministic dynamics.

The reward collection problem. Here we consider an animal who interacts with an environment over a sequence of states, actions and rewards. The animal’s goal is to select actions in a way that maximizes the total reward collected. How should an animal do this?

When the environment is unknown, to maximize reward an animal must trade-off between exploration, and exploitation. The standard way of approaching this tradeoff is to assume,

The agent must try a variety of actions and progres-

Significance Statement

An animal exploring a new environment quickly finds itself with a difficult question: Is it better to keep exploring, or exploit and stick with the known? This problem, the explore-exploit dilemma, is mathematically intractable but common to all animals. When one mathematical problem can’t be solved, it’s often good to find another related problem that can be. In this work we show that when exploration is re-imagined as an open-ended search to learn as much as possible, curiosity, it becomes possible to find simple tractable solutions for all explore-exploit problems, if that is curiosity is just as important as reward for long-term survival.

The authors have no conflicts of interest to declare.

¹To whom correspondence should be addressed. E-mail: Erik.Exists@gmail.com

sively favor those that appear to be best” (?).

(Later on we will suggest two fairly radical changes to this standard formula. Don’t stop searching progressively, do so suddenly. Explore only for learning sake, out of curiosity (?).)

Markovian formalities. Our Markov decision process (MDP) consists of states of real valued vectors \mathbf{S} from a finite set \mathcal{S} of size n . As are actions, \mathbf{A} from a finite set \mathcal{A} of size k . Rewards are non-negative real numbers, R . Transitions from state to state are handled by the transition function $\Lambda(\mathbf{S}, \mathbf{A})$. In our mathematics we assume Λ is a deterministic function of its inputs. In our simulations we assume it is a stochastic function.

Bellman optimal reward collection. These preliminaries in place we can write down the value function definition, and the standard Bellman equation (?) for reward collection, as given by reinforcement learning theory (?). For simplicity we assume below that the horizon is finite, bounded by T .

$$\begin{aligned} V_R^*(\mathbf{S}) &= \operatorname{argmax}_{\mathbf{A} \in \mathcal{A}} \left[\sum_{k=0}^T R \right] \\ &= \operatorname{argmax}_{\mathbf{A} \in \mathcal{A}} \left[R_t + \sum_{k=0}^T R_{t+1} \mid \mathbf{S}_t = \mathbf{S}, \mathbf{A}_t = \mathbf{A} \right] \quad [1] \\ &= \operatorname{argmax}_{\mathbf{A} \in \mathcal{A}} \left[R_t + V_R^*(\mathbf{S}_{t+1}) \mid \mathbf{S}_t = \mathbf{S}, \mathbf{A}_t = \mathbf{A} \right] \\ &= R_0 + V_R^*(\mathbf{S}_t) + V_R^*(\mathbf{S}_{t+1}), \dots \end{aligned}$$

The second to last equation is the optimal Bellman equation for V_R^* . This can be restated more compactly as,

$$V_R^*(\mathbf{S}) = \operatorname{argmax}_{\mathbf{A} \in \mathcal{A}} \left[R_t + V_R^*(\Lambda(\mathbf{S}, \mathbf{A})) \right] \quad [2]$$

Formal regret. Regret minimization is a standard metric of optimality in reinforcement learning (?). Formally, given *any* optimal value solution we can define regret G in terms of this, given by $G = V^* - V_t$, where V_t is the value of the chosen action \mathbf{A}_t . An optimal value algorithm can then be restated, as a “best” series of best actions ($\mathbf{A}_t, \mathbf{A}_{t+1}, \mathbf{A}_{t+2}, \dots$) for which $\sum_{k=0}^T G_t = 0$.

The information collection problem. Imagine we wish to arrive at an equation to maximize curious behavior, in terms of information gained, or information value. Assume, as a stand-in, \hat{E} represents the information value of some observation. What we would like is a solution for this term \hat{E} that would let us arrive at an equation like Eq. 2, but for information value in place of reward. This could hypothetical equation could be written as,

$$V_E^*(\mathbf{S}) = \operatorname{argmax}_{\mathbf{A} \in \mathcal{A}} \left[\hat{E}_t + V_E^*(\Lambda(\mathbf{S}, \mathbf{A})) \right] \quad [3]$$

How can we arrive at such an equation that is both general, and therefore suitable for use across species, and how can we do so when we necessarily limit ourselves to non-Markovian kinds of memory?

non-Markovian formalities. Observations. Animals at different stages in development, and in different environments, express curious behavior about the environment, their own actions, and their own mental processes. We assume a parsimonious way to handle this is be assuming curious animals often “bind” (?) the elements of environment, the Markov process ($\mathbf{S}, \mathbf{A}, \mathbf{S}', R$), into single observations \mathbf{X} . For convenience, we assume observations are embedded in a finite real space, $\mathbf{X} \in \mathcal{X}$ of size l . This binding is imagined to happen via a communication channel, T , identical to our observation function, $T(\mathbf{S}, \mathbf{A}, \mathbf{S}', R, \mathbf{M})$. (\mathbf{M} is the animal’s memory, that we define below.)

Observations made by T can be internally generated from memory \mathbf{M} , or externally generated from the environment (\mathbf{S}, \mathbf{S}'). Or both in combination, as happens in recurrent neural circuits. Observations may contain action information \mathbf{A} , but this is optional. Observations may contain reward information R , but this optional. Though omitting reward would probably hamper solving explore-exploit problems.

Memory. Sustained firing, the strength between two synapses, elevated Calcium concentration are three simple examples of learning and memory in the nervous system. Each of these examples though can be represented as a vector space. In fact, most any memory system can be so defined.

We define memory as a set of real valued numbers, embedded in a finite space that is closed and bounded with p dimensions. A learning function is then any function f that maps observations \mathbf{X} into memory \mathbf{M} . This f is considered to be a valid learning function as long as the mapping changes \mathbf{M} some small amount. A non-constant mapping, in other words. Using recursive notation, this kind of learning is denoted by $\mathbf{M} \leftarrow f(\mathbf{X}, \mathbf{M})$. Though we will sometimes use \mathbf{M}' to denote the updated memory in place of the recursion. Other times we will add subscripts, like $(\mathbf{M}_t, \mathbf{M}_{t+1}, \dots)$, to express the path of a memory.

We do not assume memory or the learning function are Markovian.

We can also define a forgetting function that *can be any function* that inverts the memory, $f^{-1}(\mathbf{X}; \mathbf{M}') \rightarrow \mathbf{M}$. Forgetting might not seem crucial, but it will prove to be so later on in developing exploration algorithms that provably maximize \hat{E} .

For individual animals we assume f and f^{-1} have been chosen by evolution and experience to be learnable, efficient, and have sufficiently useful inductive bias for their particular environment (?). This assumption is not necessary for the mathematical results which follow.

Axiomatic information value. To make our eventual union proper, we first need to separate reward value from information value in a principled and general way. We do this axiomatically. We reason instead that the value of any observation made by an animal depends entirely on what an animal learns by making that observation—literally how it changes that animal’s memory. This view let us sidestep a problem present in most working models of curiosity. They tend to conflate the learning rule, with curiosity itself (?).

Our definition of information value closely follows the motivations of Shannon, and his definition of a communication channel. If we have a memory \mathbf{M} , which has been learned by f over a history of observations, $(\mathbf{X}_0, \mathbf{X}_1, \dots)$, Can we measure how much value the next observation \mathbf{X}_t should have? We

think this measure \hat{E} should have certain intuitive properties.

Axiom 1 (Axiom of Memory). \hat{E} depends only on the difference $\delta\mathbf{M}$ between \mathbf{M} and \mathbf{M}' .

That is, the value of an observation \mathbf{X} depends only on how the memory changes, by f .

Axiom 2 (Axiom of Novelty). $\hat{E} = 0$ if and only if $\delta\mathbf{M} = 0$.

That is, an observation that doesn't change the memory has no value.

Axiom 3 (Axiom of Scholarship). $\hat{E} \geq 0$.

That is, all (new) information is in principle valuable even if its consequences are later found to be negative.

Axiom 4 (Axiom of Completeness). \hat{E} should increase monotonically with the total change in memory.

That is, there should be a one-to-one relationship between how memory changes and how information is valued.

Axiom 5 (Axiom of Equilibrium). For the same observation \hat{E} should approach 0 in finite time.

That is, learning on \mathbf{M} makes continual progress toward equilibrium (i.e. self-consistency) with each observation, after a finite time period has passed. After this it will approach a steady-state value of $\delta\mathbf{M} \rightarrow 0$ and $\hat{E} \rightarrow 0$.

All we have really done in these axioms is remove any mention of a learning rule, its properties, or learning target, its properties from previous definitions (?????). This is important not for its own sake but because with these axioms we can consider information value, and therefore curiosity, independent of any learning semantics, or any semantics at all. Anything we prove in this more abstract setting will therefore be true for any learning algorithm, who meets our requirements.

Put another way, if the communication channel is agnostic to semantic considerations, as Shannon's requires, we reason information value should be agnostic as well.

Practical information value. Meeting our requirements is not difficult. The practical example of axiomatic value we use throughout this paper is based on the geometric norm, $\|\cdot\|$ (as shown in Fig. 2). Though not necessarily a unique solution, a norm on the memory gradient is a simple, computational convenience, way to satisfy Axioms 1-4 (Eq 4).

$$\hat{E} = \|\nabla\mathbf{M}\| \quad [4]$$

To satisfy Axiom 5 we further require memory dynamics eventually decelerate, $\nabla^2\mathbf{M} < 0$ for all $t \geq T^*$. Informally, the idea here is that any notion of sensible and useful learning must converge. We take this to mean that in finite time bound, T^* , \hat{E} must approach 0 (as shown in Fig. 2). We term this consistent learning.

In practice here we will not work with the gradient, but will use a discrete time map in its place,

$$\hat{E} \approx \|f(\mathbf{M}, \mathbf{X}) - \mathbf{M}\| \quad [5]$$




Fig. 2. Practical (geometric) information value, and our constraints on its dynamics. **a.** This panel illustrates a two dimensional memory. The information value of two observations \mathbf{X}_1 and \mathbf{X}_2 depends on the norm of memory with learning. Here we show this distance as a euclidean norm, denoted as a black arrow. **b-c** This panel illustrates learning dynamics with time (over a series of observations that are not shown). If information value becomes decelerating in finite time bound, then we say that learning is consistent with Def. 2. This is shown in panel b. The bound is depicted as a dotted line. If learning does not decelerate, then it is said to be inconsistent (Panel c). *It is important to note:* our account of information value and curiosity does not work when learning is inconsistent.

Bellman optimal information collection. The common way to arrive at Bellman solution is to assume a Markov decision space. This is a problem for our definition of memory, which we assume has a natural long-term path dependence. This breaks the Markov assumption, and it's needed claim of optimal substructure. To get around this we proved that exact forgetting of the last observation is another way to achieve optimal substructure, and so to derive a Bellman solution.

In the theorem below which shows the optimal substructure of \mathbf{M} we implicitly assume that $(X = \mathbf{S}, \mathbf{A}, f, \text{ and } \Lambda)$ are given, and that Λ is deterministic,

[Optimal substructure] If $V_{\pi_{\hat{E}}}^*$ is the optimal information value given by policy $\pi_{\hat{E}}$, a memory \mathbf{M}_t has optimal substructure if the last observation X can be removed from \mathbf{M} , by $\mathbf{M} - \mathbf{1}_t = f^{-1}(\mathbf{X}, \mathbf{M}_t)$ such that the resulting value $V_{t-1}^* = V_t^* - E_t$ is also optimal.

Given an arbitrary starting value $E_0 > 0$, the best solution to maximizing \hat{E} is can be given by a Bellman equation (Eq. 6).

$$V_{\hat{E}}^*(\mathbf{S}) = \operatorname{argmax}_{\mathbf{A} \in \mathbf{A}} [\hat{E}_t + V_{\hat{E}}^*(\Lambda(\mathbf{S}, \mathbf{A}))] \quad [6]$$

As long as learning about the different observations \mathbf{X} is independent, and as long as learning continues until steady-state, defined as $E_t < \eta$, there are many equally good solutions to the information collection problem above. Under these conditions we can therefore simplify Eq. 6 further, giving Eq. 7.

$$V_{\hat{E}}^*(\mathbf{X}) = \operatorname{argmax}_{\mathbf{A}} [E_t + E_{t+1}] \quad [7]$$

A hypothesis about boredom. The central failure mode of curiosity is what we'll call minutia, defined as those learning that has little to no use, ever. A good example of this is to, "Imagine a scenario where the agent is observing the movement of tree leaves in a breeze. Since it is inherently hard to model breeze, it is even harder to predict the location of each leaf" (?). This will imply, they note, a curious agent will always remain

curious about chaos in the leaves even though they have no hope of successful learning of them.

To limit curiosity we introduce a penalty term, η which we treat as synonymous with boredom. We consider boredom an adaptive trait, in other words. Others have considered boredom arguing it is a useful way to motivate aversive tasks (?), or curiosity (?). We take a slightly different view. We treat boredom as a tunable free parameter, $\eta \geq 0$ and a means to ignore marginal value by requiring curious exploration to stop once $E \leq \eta$.

Optimal exploration. Having a policy π_E^* that optimally maximizes \hat{E} does not necessarily ensure that exploration is of good quality. We think it is reasonable to call an exploration good if it meets the criteria below.

1. Exploration should visit all available states of the environment at least once.
2. Exploration should cease when learning has plateaued.
3. Exploration should take as few steps as possible to achieve 1 and 2.

In Theorems 1 and 2 we prove that our Bellman solution π_E satisfies these, when $\eta > 0$ and when the observations \mathbf{X} contain complete state information, $\mathbf{S} \in \mathbf{X}$. See the Appendix for the proofs proper.

Theorem 1 (Complete exploration). *Given some arbitrary value E_0 , an exploration policy governed by π_E^* will visit all states $\mathbf{S} \in \mathbf{S}$ in a finite number of steps T .*

Theorem 2 (Efficient exploration). *An exploration policy governed by π_E^* will only revisit states for where information value is marginally useful, defined by $\hat{E} > \eta$.*

A hypothesis of equal importance. Searching in an open-ended way for information must be less efficient than a direct search for reward? In answer to this question, we make two conjectures.

Conjecture 1 (A hypothesis of equal importance). *Reward value and information value are equally important for survival, in changing environments*

It's worth reiterating here, that curiosity is a primary drive in most, if not all, animals (?). It is as strong, if not sometimes stronger, than the drive for reward (???).

If both reward and information are equally important, then time is not wasted in a curious search, and so the process is not necessarily inefficient or even indirect. Instead, we offer a new question. Is curiosity practical enough to serve as a useful "trick" to solve reward collection search problems? We answer this with a second conjecture,

Conjecture 2 (The curiosity trick). *When learning is possible, curiosity is a sufficient solution for all exploration problems.*

It's worth noting here that curiosity, as an algorithm in machine learning, is highly effective at solving broad kinds of optimization problems (????????).

The problem of mixing reward and information collection. If reward and information are equally important, and curiosity is sufficient, how can an animal go about balancing them? If you accept our conjectures, then answering this question optimally is the same as solving the dilemma. To solve this dual value learning problem we've posed—information and reward maximization—we need an algorithm that can maximize the value of a mixed sequence, $V_{\hat{E}R}$. For example,

$$V_{\hat{E}R}^*(\mathbf{S}) = \operatorname{argmax}_{\mathbf{A} \in \mathbf{A}} \left[\sum_{k=0}^T \hat{E}_t + \hat{E}_{t+1} + R_{t+2} + \hat{E}_{t+3} + R_{t+4} + \dots \right] \quad [8]$$

A win-stay, lose-switch solution. A similar dual value optimization problem to Eq. 8 was posed, solved, with its regret bounded, by Robbins (?) in his derivation of the win-stay lose-switch rule. This WSLS approach has since proven useful in decision making and reinforcement learning (??), Bayesian sampling (?), and especially in game theory (?). We also put it to use here.

We have come to think about our two optimal value policies π_R and $\pi_{\hat{E}}$ as two "players" in game for behavioral control of an animal's moment-by-moment actions (?). The simplest solution to this game, in terms of computational complexity, is a WSLS rule. We denote this rule as Π_π , and define it here by the inequalities below.

$$\Pi_\pi = \begin{cases} \pi_{\hat{E}}^* & : \hat{E} - \eta > R + \rho \\ \pi_R & : \hat{E} - \eta < R + \rho \end{cases} \quad [9]$$

We assume no matter which policy is in control in Eq. 9, each policy can learn from (be updated using) the other policies actions, and observations. That is, while the two policies compete for control, they learn from each other in a cooperative fashion.

Here η is the boredom term we defined earlier. We also introduce the reward bias term ρ , where $\rho > \eta$. This bias ensures no ties can occur, and that as \hat{E} decreases to 0, as it must, the default policy is reward collection. That is, by π_R .

In Eq. 9 we compare reward R and information value \hat{E} but have not made sure they have the same "units", or are otherwise comparable. We sidestep this issue by limiting the probability of having zero reward to be, itself, non-zero. That is, $p(R_t = 0) > 0$. We also limit E_0 , the first set of values, to be positive and non-zero when used with boredom, $(E_0 - \eta) > 0$. These constraints ensure complete "good" exploration (defined above), which in turn ensures optimal reward learning is possible irrespective of the relative scales of R and \hat{E} .

To provide some intuition for Eq. 9, imagine that in the previous action an animal observed, arbitrarily, 0.7 units of information value, \hat{E} , and no reward (i.e. $R = 0$). Using our rule on the next action the animal should then choose its exploration policy $\pi_{\hat{E}}^*$. If it explored last time, this is a "stay" action. If not, it is a "switch". Let's then say that with this next observation, \hat{E} decreases to 0.2 and a reward value of 1.0 is observed. Now the animals should switch its strategy for the next round, to exploit instead.

In general, if there was more reward value than information value last round ($R > \hat{E}$), our rule dictates an animal should choose the reward policy. If information value dominated, $R < \hat{E}$, then it should choose the information gathering policy.

The WSL rule in Eq 9 when applied to exploration-exploitation problems can guarantee three things, shown below.

Theorem 3 (No regret - reward value). *When π_R is in control under Π_π , all actions are zero regret in terms of V_R . That is, $\sum_{k=0}^T G = 0$.*

Theorem 4 (No regret - information value). *When $\pi_{\hat{E}}$ is in control under Π_π , all actions are zero regret in terms of $V_{\hat{E}}$. That is, $\sum_{k=0}^T G = 0$.*

Theorem 5 (No regret - mixed values). *When either $\pi_{\hat{E}}$ or π_R is in control under Π_π , all actions are zero regret in terms of $V_{\hat{E}R}$. That is, $\sum_{k=0}^T G = 0$.*

In each of these Theorems we assume our definitions for \mathbf{X} , \mathbf{A} , \mathbf{M} , f , Λ , and Π_π are implicitly given, and a finite horizon T . Proofs for Theorems 4 and 4 follow from both reinforcement learning and our definition of information optimization having a Bellman optimality. The proof for the optimality of $V_{\hat{E}R}$ (Th 4) is provided in the Appendix. The intuition for this proof is simple that a greedy algorithm made over two Bellman solutions is itself a Bellman solution.

In Eq. 9 we compare reward R and information value \hat{E} but have not ensured they have the same “units”, or are comparable valuables. We sidestep this issue by limiting the probability of having zero reward to be, itself, non-zero. That is, $p(R_t = 0) > 0$. We also limit E_0 , the first set of values, to be positive and non-zero when used with boredom, $(E_0 - \eta) \geq 0$. These constraints ensure complete “good” exploration (defined above) which in turn ensures optimal reward learning is possible.

Homeostatic adjustments. We have been studying cases where reward value is fixed. That is, where the environment reward R is equal to its subjective value to the animal, \hat{R} . In natural behavior though the motivational value of reward often declines as the animal reaches satiety (? ? ?). This is called reward homeostasis. Fortunately, it has been shown a reward collection policy designed for one style will work for the other (?). However to use a form similar to Eq. 9 for the dilemma under homeostatic reward value, we need to make one modification. This is shown in Eq. 10.

$$\Pi_\pi = \begin{cases} \pi_{\hat{E}}^* & : \hat{E} - \eta > R - \rho \\ \pi_R & : \hat{E} - \eta < R - \rho \end{cases} \quad [10]$$

Under homeostasis ties between values in Eq. 9 should be broken in favor of exploration, because if exploitation remains the default, reward homeostasis is not possible.

Experiments. Does our theory work in practice? Our approach does not mathematically ensure a mixed value solution to explore-exploit problems will always maximizes reward value, compared with other well-established methods. The remainder of this paper is devoted to studying our approaches performance in simulated bandit tasks. Each task was designed to either test exploration performance in a way that matches recent experimental studies, or to test the limits of curiosity.




Fig. 3. A simple four choice information foraging task. The information in this task is a yellow or blue stimulus, which can change from trial to trial. A good learner in this task is one who tries to learn the probabilities of all the symbols in each of the four choices. The more random the stimuli are in a choice, the more potential information/entropy there is to learn. Note: the information payout structure is shown below (Fig. 6a).

Information collection. The work so far has built up the idea that the most valuable, and most efficient, curious search will come from a deterministic algorithm. That is, every step strictly maximizes \hat{E} . It is this determinism which will let us resolve the dilemma, later on. A deterministic view of exploration seems at odds with how the problem is generally viewed today, which involves searching with some amount of randomness. Whereas if our analysis is correct, randomness is not needed or desirable, as it must lead to less value and a longer search.

We confirm and illustrate this result using an example of a simple information foraging task (Task 1; Fig.). This variation of the bandit task (?) replaces rewards with information, in this case colors. On each selection, the agent sees one of two colors according to a specific probability shown in Fig. 6a. When the relative color probabilities are more similar that arm has more entropy and so more to learn and more information value. Arms that are certain to present only one color lose their informative value quickly.

The results of this simple information seeking task are shown in Figure 4. Deterministic curiosity in this task generated more information value, in less time, and with zero regret when compared against a stochastic agent using more directed random search. As our work here predicts, noise only made the search happen slower and lead to regret. Besides offering a concrete example, performance in Task 1 is a first step in showing that even though π_E 's optimality is proven for a deterministic environment, it can perform well in other settings.

Reward collection. Our approach does not mathematically ensure a dual value solution to the dilemma maximizes reward value compared with other well-established methods. To find out we measured total reward collected over 7 tasks and 10 agents, including ours. Each agent's parameters were independently optimized for each task. For a brief description of each agent, see Table 1. They all have in common that their central goal is to maximize total reward value, though this is often supplemented by some other goal, an intrinsic reward or bonus (? ?).

The general form of the 7 tasks are depicted in Fig 5. As with our first task (Fig.) they are all variations of the classic multi-armed bandit (?). The payouts for each of the tasks are shown separately in Fig. 6. Every trial has a set of n choices. Each choice returns a “payout”, according to a predetermined probability. Payouts are information, a reward, or both. Note

Deterministic versus stochastic curiosity in a random world

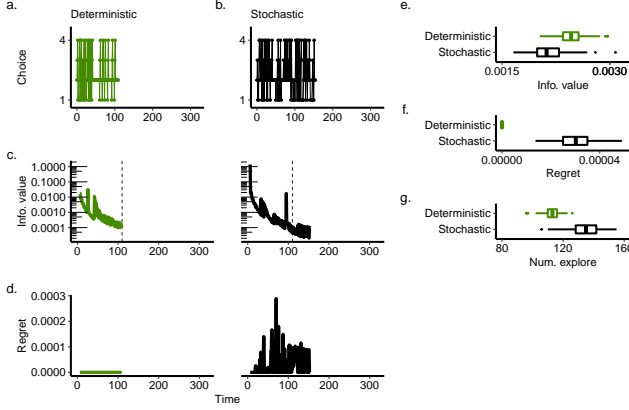


Fig. 4. Comparing deterministic versus stochastic variations of the same curiosity algorithm, in a simple information foraging task (Task 1). Deterministic results are shown in the left column, and stochastic are shown in the right. The hyperparameters for both models were found by random search, which is described in the Methods. **a-b.** Examples of choice behavior. **c-d.** Information value plotted with time for the behavior shown in a-b. **e-f.** Regret plotted with time for the behavior shown in a-b. Note how our only deterministic curiosity generates zero regret, inline with theoretical predictions. **e.** Average information value for 100 independent simulations. Large values mean a more efficient search. **f.** Average regret for 100 independent simulations. Ideal exploration should have no regret. **g.** Number of steps it took to reach the boredom threshold *etc.* Smaller values imply a faster search.




Fig. 5. A 4 choice reward collection task. The reward is numerical value, here a 1 or 0. A good learner in this task is one who collects the most rewards. The task depicted here matches that in Fig 6b. Every other reward collection task we study has the same basic form, only the number of choices increases and the reward spaces are more complex.

that, as in Task 1, information was symbolic, denoted by a color code, “yellow” and “blue” and as is custom reward was a positive real number.

The results in full for all tasks and exploration strategies are shown in Fig. . All agents, including ours, used the same exploitation policy based on the temporal difference learning rule () (Methods).

Task 1. is an information gathering task, which we discussed already above. *Task 2* was designed to examine reward collection, in probabilistic reward setting very common in the decision making literature (). Rewards were 0 or 1. The best choice had a payout of $p(R = 1) = 0.8$. This is a much higher average payout than the others ($p(R = 1) = 0.2$). At no point does the task generate symbolic information. See, Fig. ??b.

The results for Task 1 were discussed above. As for Task 2, we expected all the exploration strategies to succeed. While this was indeed the case, our deterministic curiosity was the top-performer in terms of median rewards collected, though

Table 1. Exploration strategies.

Name	Class	Exploration strategy
Curiosity	Deterministic	Maximize information value
Random/Greedy	Random	Alternates between random exploration and greedy with probability ϵ .
Decay/Greedy	Random	The ϵ parameter decays with a half-life τ
Random	Random	Pure random exploration
Reward	Extrinsic	Softmax sampling of reward value
Bayesian	Extrinsic + Intrinsic	Sampling of reward value + information value
Novelty	Extrinsic + Intrinsic	Sampling of reward value + novelty signal
Entropy	Extrinsic + Intrinsic	Sampling of reward value + action entropy
Count (EB)	Extrinsic + Intrinsic	Sampling of reward value + visit counts
Count (UCB)	Extrinsic + Intrinsic	Sampling of reward value + visit counts

by a small margin (Fig a).

Task 3. was designed with very sparse rewards () and there were 10 choices, making this a more difficult task (Fig. ??c). Sparse rewards are a common problem in the natural world, where an animal may have to travel and search between each meal This is a difficult but not impossible task for vertebrates () and invertebrates (). That being said, most reinforcement learning algorithms will struggle in this setting because the thing they need to learn, that is rewards, are often absent. In this task we saw quite a bit more variation in performance, with the novelty-bonus strategy taking the top slot (this is the only time it does so).

Task 4. was designed with deceptive rewards. By deceptive we mean that the best long-term option presents itself initially with a decreasing reward value (Fig. ??d). Such small deceptions abound in many natural contexts where one must often make short-term sacrifices (). It is well known that classic reinforcement learning will often struggle in this kind of task (). Here our deterministic curiosity is the only strategy that reaches above chance performance. Median performance of all other strategies are similar to the random control (Fig c).

Task 5 was designed to fool curiosity, our algorithm, by presenting information that was utterly irrelevant to reward collection, but had very high entropy and so “interesting” to our algorithm. We fully anticipated that this context would fool just our algorithm, with all other strategies performing well since they are largely agnostic to entropy. However, despite being designed to fail, deterministic curiosity still produced a competitive performance to the other strategies (Fig d).

Tasks 6-7 were designed as a pair, with both having 121 choices, and a complex payout structure. Tasks of this size are at the limit of human performance (). We first trained all learners on *Task 6*, then tested them in *Task 7* which identical to 6, except the best payout arm is reset to be worst (Fig. ??e-f). In other words Tasks 6 and 7 were joined to measure learning in a high dimensional, but shifting, environments.

In Task 6 deterministic curiosity performed well, securing a second place finish. We note the Bayesian strategy outperformed our approach. However, under the sudden non-stationarity in reward when switching tasks, the top Bayesian model became the worst on Task 7, and deterministic curiosity took the top spot. Compare Fig. e to f. This is the robustness that we’d expect for any curiosity algorithm, whose main goal is to learn everything unbiased by other objectives. The environment and the objectives do change in real life, and so

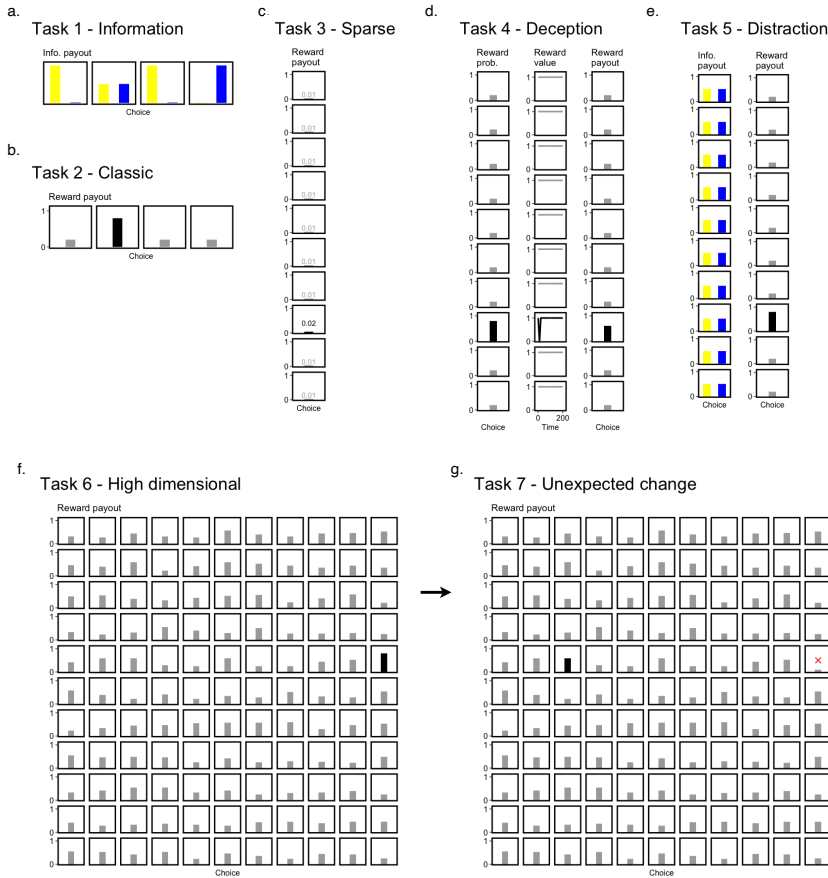


Fig. 6. Payouts for Tasks 1 - 7. Payouts can be information, reward, or both. For comments on general task design, see Fig 5. **a.** A classic four-choice design for information collection. A good learner should visit each arm, but quickly discover that only arm two is information bearing. **b.** A classic four-choice design for testing exploration under reward collection. The learner is presented with four choices and it must discover which choice yields the highest average reward. In this task that is Choice 2. **c.** A ten choice sparse reward task. The learner is presented with four choices and it must discover which choice yields the highest average reward. In this task that is Choice 8 but the very low overall rate of rewards makes this difficult to discover. Solving this task with consistency means consistent exploration. **d.** A ten choice deceptive reward task. The learner is presented with 10 choices, but the choice which is the best on the long-term (>30 trials) has a lower value in the short term. This value first declines, then rises (see column 2). **e.** A ten choice information distraction task. The learner is presented with both information and rewards. A good learner though will realize the information does not predict reward collection, and so will ignore it. **f.** A 121 choice task with a complex payout structure. This task is thought to be at the limit of human performance. A good learner will eventually discover choice number 57 has the highest payout. **g.** This task is identical to **a.**, except for the high payout choice being changed to be the lowest possible payout. This task tests how well different exploration strategies adjust to simple but sudden change in the environment.

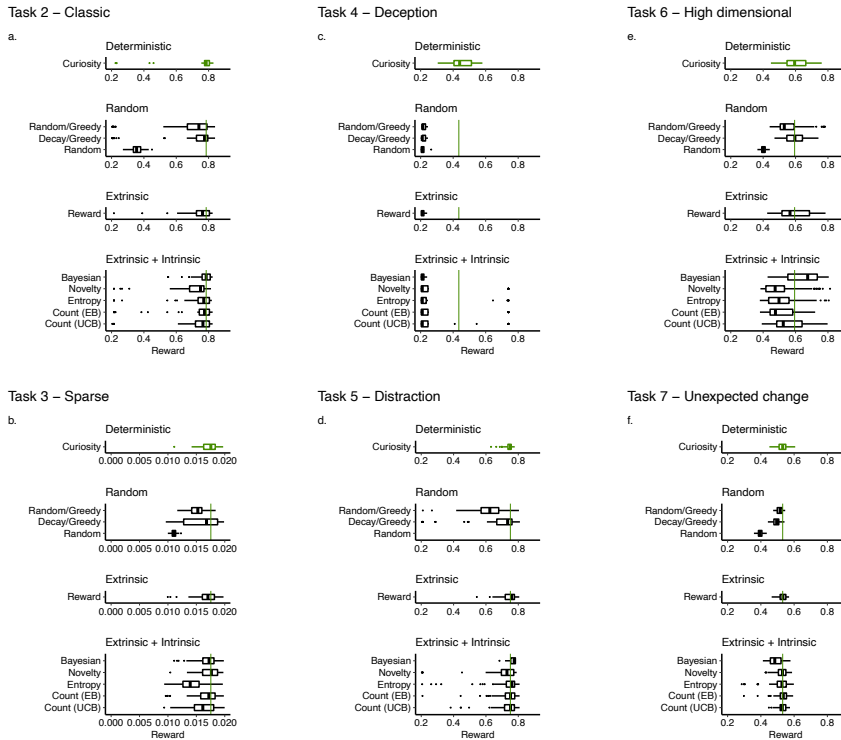


Fig. 7. Summary of reward collection (Tasks 2-7). The strategies in each panel are grouped according to the class of search they employed (Curiosity. Random, Extrinsic reward or Extrinsic + Intrinsic rewards). **a.** Results for Task 2, which has four choices and one clear best choice. **b.** Results for Task 3, which has 10 choices and very sparse positive returns. **c.** Results for Task 4, whose best choice is initially “deceptive” in that it returns suboptimal reward value over the first 20 trials. **d.** Results for Task 5, which blends the information foraging task 1 with a larger version of Task 2. The yellow/blue stimuli are a max entropy distraction which do not predict the reward payout of each arm. **e.** Results for Task 6, which has 121 choices and a quite heterogeneous set of payouts but still with one best choice. **f.** Results for Task 7, which is identical to Task 6 except the best choice was changed to be the worst. The learners from Task 6 were trained on this Task beginning with the learned values from their prior experience – a test of robustness to sudden change in the environment. *Note:* To accommodate the fact that different tasks were run different numbers of trials, we normalized total reward by trial number. This lets us plot reward collection results for all tasks on the same scale.

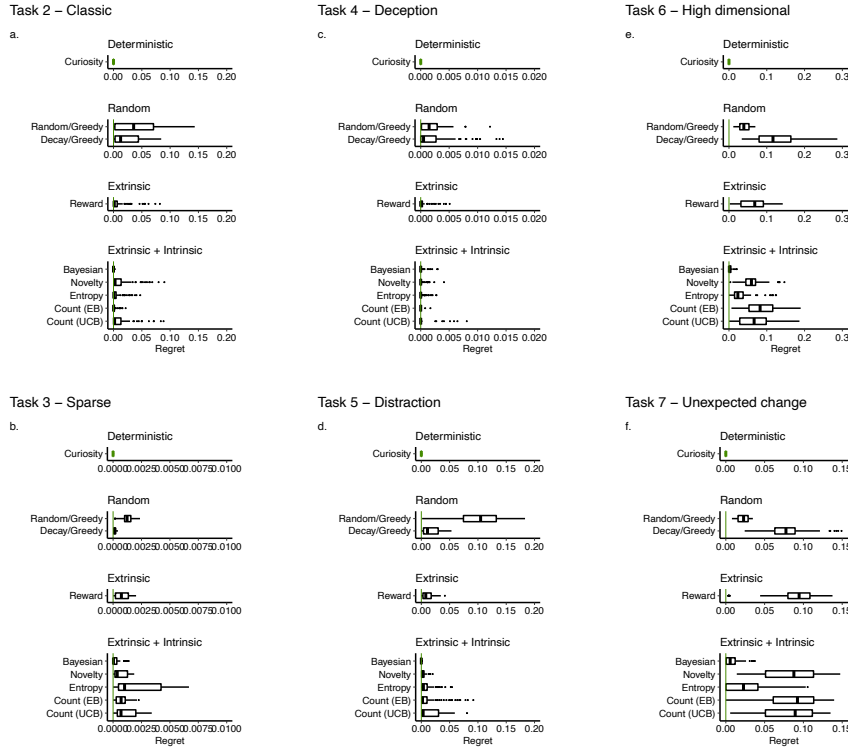


Fig. 8. Summary of total regret (Tasks 2-7). See the previous figure for details.

Exploration strategies Ten 'animals'

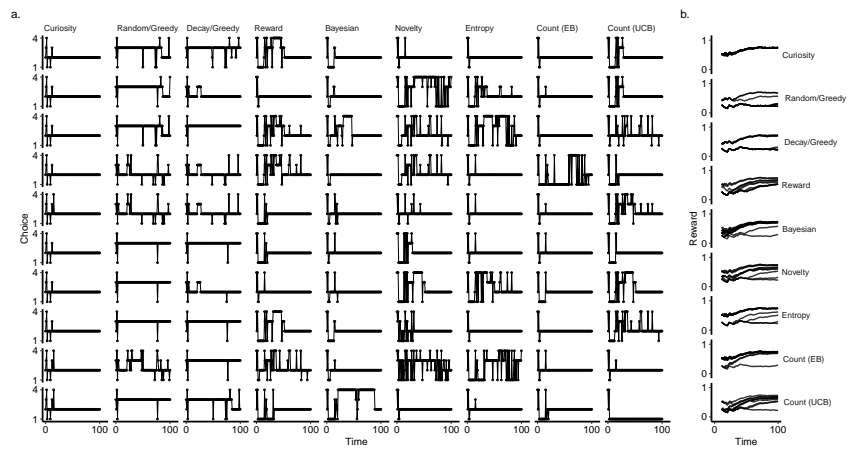


Fig. 9. Examples of exploration strategies (Task 1). Here we simulated different sets of randomly chosen hyperparameters on the same task, for all agents. We take the view here that each parameter choice acts akin to a unique "animal" (?). So this figure estimates, to some degree, the degree and kind of variability we might expect in a natural experiment.

we must be prepared for this. Note how the other less-biased-toward-reward-value exploration models (count-based, novelty, and entropy models) also saw gains, to a lesser degree.

Robustness. Unlike idealised simulations, animals cannot pre-optimize their search parameters for every task. We have so far been using carefully chosen parameters. For example, a good choice of boredom η is critical for our definition of curiosity, and was optimized carefully. The next question we considered, albeit somewhat qualitatively, is – how much does the choice of η matter?

When we move from an optimized choice of η , to a value about 3 fold smaller, we see optimized value converged far more quickly (Fig. 10a-b), generating more reward over time (Fig. 10c-d)

The next question we considered is, how sensitive is this parameter compared to random parameter choices for all other nine agents? As an answer to this question we therefore explored reward collection as a function of 1000 randomly chosen model parameters, and reexamined performance on Tasks 1 and 7. These two were chosen to represent an “easy” task, and a “hard” one, in terms of search size. The results are shown in Fig. 11, where we plot ranked performance on random parameters.

No matter the parameter choices, and as we observed in our other model comparisons, exploitation by curiosity produced top-3 performance on both tasks (Fig. 11a-bd. Most interesting to note is the performance for the worst 10 random parameters. Here curiosity was markedly better, with substantially less variability across different parameter schemes.

We can explain this last result in the following way. All the other exploration strategies use parameters to tune the degree of exploration (i.e. the level of noise). With our approach to curiosity, however, we tune only when exploration should stop (by η). That is exploration is pure, and unbiased by reward or noise level. And so it seems, in these examples anyway, it is far more robust to tune the stopping point for pure exploration, rather than the degree of exploration.

The final control questions we considered were, how does our (deterministic) approach to curiosity decline under stochastic sampling? And, does it decline more or less quickly than other agents? To induce stochastic behavior in our model we sampled its value estimates via a Softmax; the same mechanism used to drive stochastic search in several other of the alternatives. Results for this are shown in Figure 12.

Adding noise to our approach eventually degrades performance. But performance falls off no more quickly than several of the other models. Contrariwise, when stochasticity becomes low, most the alternatives show substantial declines or variability in their performance.

Discussion

We have offered a way around the exploration-exploitation dilemma. We proved the classic dilemma, found when optimizing exploration for reward value, can be resolved by exploring instead for curiosity. And that in this form the rational trade-off we are left with is solved with a classic strategy taken from game theory. The win-stay lose-shift strategy, that is.

This is a controversial claim. So we will address some of its criticisms as a series of questions. The final question explains how this approach can be tested, with behavioral experiments.

Why curiosity? Our use of curiosity rests on three well established facts. 1. Curiosity is a primary drive in most, if not all, animal behavior (?). 2. It is as strong, if not sometimes stronger, than the drive for reward (? ? ?). 3. Curiosity as an algorithm is highly effective at solving difficult optimization problems (? ? ? ? ? ? ? ? ?).

What about information theory? In short, the problems of communication and value are different problems that require different theories.

Weaver (?) in his classic introduction to the topic, describes information as a communication problem with three levels. a. The technical problem of transmitting accurately. b. The semantic problem of meaning. c. The effectiveness problem of changing behavior. He then describes how Shannon’s work addresses only problem A. It has turned out that Shannon’s separation of the problem allowed the information theory to have an incredibly broad application.

Valuing information is not, at its core, a communications problem. Consider the very simple example Bob and Alice, having a phone conversation and then Tom and Bob having the exact same conversation. The personal history of Bob and Alice will determine what Bob learns in their phone call, and so determines what he values in that phone call. These might be completely different from what he learns when talking to Tom, even if the conversations are identical. What we mean by this example is that the personal history defines what is valuable on a channel, and this is independent of the channel and the messages sent. We have summarized this diagrammatically, in Fig. 13.

There is, we argue, an analogous set of levels for information value as those Weaver describes. a. There is the technical problem of judging how much was learned. b. There is the semantic problem of both what this learning “means”, and also what its consequences are. c. There is an effectiveness problem of using what was learned to some other effect. In many ways b and c are similar to those in the information theory. It is a. which is most different.

Is this too complex? Perhaps turning a single objective into two, as we have, is too complex an answer. If this is true, then it would mean our strategy is not a parsimonious solution to the dilemma. We could or should reject it on that alone?

Questions about parsimony are resolved by considering the benefits versus the costs. The benefits to curiosity-based search is that it leads to no regret solutions to exploration, to exploration-exploitation, and that it so far seems to do very well in practice, at reward collection. At the same time curiosity-as-exploration can also be building a model of the environment, useful for later planning (? ?), creativity, imagination (?), while also building, in some cases, diverse action strategies (? ? ? ?).

The cost is of curiosities indirectness, which could make it slower to converge in principle. There is also the risk of focusing on minutia, which we discuss in the main text. Our numerical simulation suggests neither of these costs are not fatal for finding practical successes.

Is this too simple? Solutions to the regular dilemma are complex and require sophisticated mathematical efforts and complex computations, when compared to our solution. Put another way, our solution may seem too simple to some. So have

Optimizing boredom for reward collection 100 simulated experiments

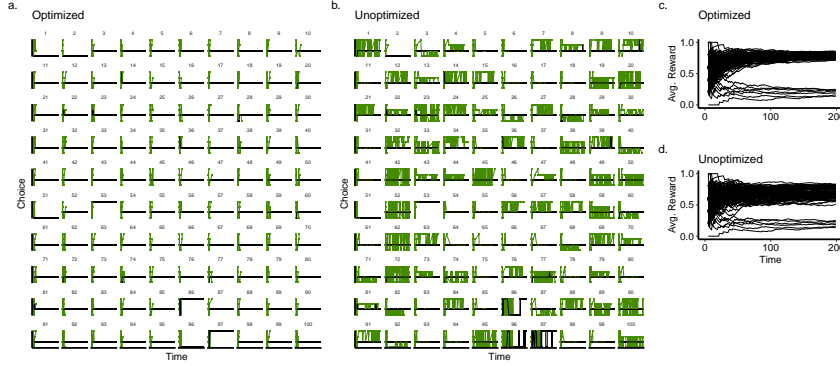
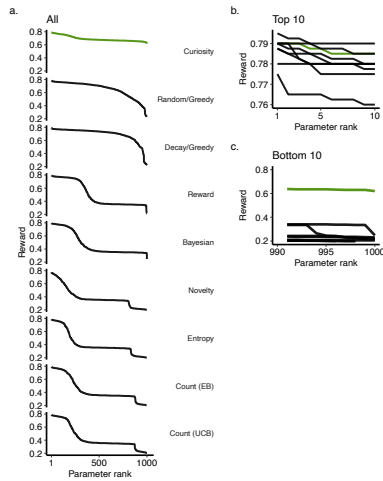


Fig. 10. The importance of boredom during reward collection (Task 2). One hundred example experiments, each began with a different random seed. **a.** Behavioral choices with optimized boredom. **b.** Behavioral choices with unoptimized boredom. **c,d.** Average reward as a function of time for optimized (c) and unoptimized (d) boredom.

Reward collection and parameter choice Task 1 – Classic



Task 6 – High dimensional

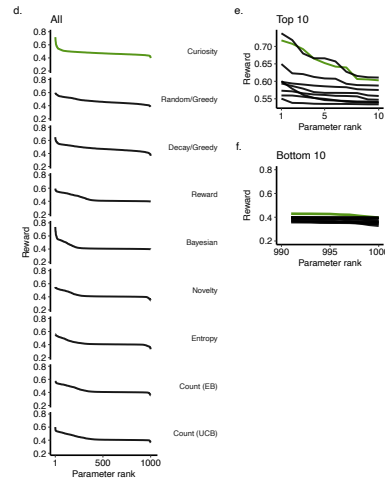


Fig. 11. Parameter sensitivity (Task 2 and 6). We wanted to evaluate how robust performance was random hyperparameters, rather than those which were carefully choose. A very robustness exploration algorithm would produce strong performance with both the best parameter choices, and the worst parameter choices. **a,d** Total reward for 1000 randomly chosen hyperparameters, for each of the agents. **b,e** Performance with the top 10 random hyperparameters, after ranking. Curiosity (green) compared to all the others. **c,f** Performance with the bottom 10 random hyperparameters, after ranking.

The effect of noise on stochastic search




Fig. 12. Noise across strategies. Adding noise to degrades performance to nearly the same degree (Task 2)




Fig. 13. The relationship between the technical problem of communication with the technical problem of value. Note how value is not derived from the channel directly. Value is derived from learning about observations from the channel, which is in turn dependent on memory and its past history.

we “cheated” by changing the problem in the way that we have?

The truth is we might have cheated, in this sense: the dilemma might have to be as hard as it has seemed. But the general case for curiosity as useful in exploration is clear, and backed up by the brute fact of its widespread presence. The question is: is curiosity so useful and so robust that it can be sufficient for all exploration with learning.

The answer to this question is empirical. If our account does well in describing and predicting animal behavior, that would be some evidence for it (? ? ? ? ? ? ? ? ? ?). If it predicts neural structures (? ?), that would be some evidence for it (we discuss this more below). If this theory proves useful in machine learning and artificial intelligence research, that would be some evidence for it (? ? ? ? ? ? ? ? ? ?).

Is this a slight of hand, theoretically? Yes. It is often useful in mathematical studies to take one problem that cannot be solved, and replace it with another related problem that can be. In this case, we swap one highly motivated animal behavior (reward seeking) for another (information seeking).

Does value as a technical problem even make sense? Having a technical definition of value, free from any meaning, might seem counterintuitive property for a value measure. We argue that it is not more or less counterintuitive than stripping information of meaning in the first place. The central question for our account of value is, will it prove useful?

What about mutual information, or truth? In other prototypical examples, information value is based on how well what is learned relates to the environment (? ? ?), that is the mutual information the internal memory and the external environment. Colloquially, one might call this “truth seeking”.

As a people we pursue information which is fictional (?), or based on analogy (?), or outright wrong (?). Conspiracy theories, disinformation, and our many more mundane errors, are far too commonplace for information value to be based on mutual information, as defined above. This does not mean that holding false beliefs cannot harm survival, but this is a second order question as far as learning goes.

So you suppose there is always positive value for learning of all fictions, disinformation, and deceptions? We must. This is our most unexpected prediction. Yet, logically, it is internally consistent with our work we believe qualitatively consistent with the behavior of humans and animals alike. The fact that humans consistently seek to learn false things calls out us to accept that learning alone is a motivation for behavior.

What about information foraging? Our work is heavily inspired by information foraging (? ?). Our motivation in developing our novel strategy was that information value should not depend on only a probabilistic reasoning, as it does in information foraging. This is for the simple reason that many of the problems animals need to learn are not probabilistic problems *from their point of view*.

Was it necessary to build a general theory for information value just to describe curiosity? No. It was an idealistic

choice that worked out. The field of curiosity studies has shown that there are many kinds of curiosity. At the extreme limit of this diversity is a notion of curiosity defined for any kind of observation, and any kind of learning. If we were able to develop a theory of value driven search at this limit, we can be sure it will apply in all possible examples of curiosity that we seem sure to discover in future. At this limit we can also be sure that the ideas will span fields, addressing the problem as it appears in computer science, machine learning, game theory, psychology, neuroscience, biology, and economics.

Doesn't your definition of regret ignore lost rewards when the curiosity policy is in control? Yes. Regret is calculated for the policy which is in control of behavior, not the policy which *could* have been in control of behavior.

Is information a reward? If reward is defined as more-or-less any quantity that motivates behavior, then our definition of information value is a reward. This last point does not mean, however, that information value and environmental rewards are interchangeable. Rewards from the environment, are a conserved resource, information is not. For example, if a rat shares a potato chip with a cage-mate, it must break the chip up leaving it less food for itself. While if a student shares an idea with a classmate, that idea is not divided up.

Furthermore, if you add reward value and information value to motivate search you can drive exploration in practical and useful ways. But this doesn't change the intractability of the dilemma proper (? ? ? ?).

In other words, it depends.

But isn't curiosity impractical? It does seem curiosity is just as likely to lead away from a needed solution as towards it. Especially so if one watches children who are the prototypical curious explorers (? ?). This is why we limit curiosity with boredom, and counter it with a competing drive for reward collecting (i.e., exploitation).

We believe there is a useful analogy between science and engineering. Science can be considered an open-ended inquiry, and engineering a purpose driven enterprise. They each have their own pursuits but they also learn from each other, often in alternating iterations (?). It is their different objectives though which make them such good long-term collaborators.

But what about curiosity on big problems? A significant drawback to curiosity, as an algorithm, is that it will struggle to deliver timely results when the problem is very large, or highly detailed. First, it is worth noting that most learning algorithms struggle with this setting (? ?), that is unless significant prior knowledge can be brought to bear (? ?) or the problem can be itself compressed (? ?). Because "size" is a widespread problem in learning, there are a number of possible solutions and because our definition of learning, memory and curiosity is so general, we can take advantage of most. This does not guarantee curiosity will work out for all problems; all learning algorithms have counterexamples (?).

But what about other forms of curiosity? We did not intend to dismiss the niceties of curiosity that others have revealed. It is that these prior divisions parcel out the idea too soon. Philosophy and psychology consider whether curiosity is internal or external, perceptual versus epistemic, or specific versus

diversive, or kinesthetic (? ? ?). Computer science tends to define it as needed, for the task at hand (? ? ? ? ?). Before creating a taxonomy of curiosity it is important to first define it, mathematically. Something we have done here.

What is boredom, besides being a tunable parameter? A more complete version of the theory would let us derive a useful or even optimal value for boredom, if given a learning problem or environment. This would also answer the question of what boredom is computationally. We cannot do this. It is the next problem we will work on, and it is very important. The central challenge is deriving boredom for any kind of learning algorithm.

Does this mean you are hypothesizing that boredom is actively tuned? Yes we are predicting exactly that.

But do people subjectively feel they can tune their boredom? From our experience, no. Perhaps it happens unconsciously? For example we seem to alter physiological learning rates without subjectively experiencing it (?).

How can an animal tune boredom in a new setting? The short answer is that one would need to guess and check. Our work in Fig 11 suggests that this can be somewhat robust.

Do you have any evidence in animal behavior and neural circuits? There is some evidence for our theory of curiosity in psychology and neuroscience, however, in these fields curiosity and reinforcement learning have largely developed as separate disciplines (? ? ?). Indeed, we have highlighted how they are separate problems, with links to different basic needs: gathering resources to maintain physiological homeostasis (? ?) and gathering information to decide what to learn and to plan for the future (? ?). Here we suggest that though they are separate problems, they are problems that can, in large part, solve one another. This insight is the central idea to our view of the explore-exploit decisions.

Yet there are hints of this independent cooperation of curiosity and reinforcement learning out there. Cisek (2019) has traced the evolution of perception, cognition, and action circuits from the Metazoan to the modern age (?). The circuits for reward exploitation and observation-driven exploration appear to have evolved separately, and act competitively, exactly the model we suggest. In particular he notes that exploration circuits in early animals were closely tied to the primary sense organs (i.e. information) and had no input from the homeostatic circuits (? ? ?). This neural separation for independent circuits has been observed in "modern" high animals, including zebrafish (?) and monkeys (? ?).

What about aversive values? We certainly do not believe this is a complete theory of rewarding decisions. For example, we do not account for any consequences of learning, or any other aversive events. Accounting for these is another important next step.

The presence of many values fits well within our notion of competition between simple (pure) policies, in a win-stay loose-shift game (?). To take other values into account we need a different notation, from the inequalities of Eq. 9. For example, if we added a fear value F to our homeostatic scheme we can move to argmax notation,

$$\Pi^\pi = [\pi_E, \pi_R, \pi_F] \quad [11]$$

$$\operatorname{argmax}_{\Pi^\pi} [E - \eta, \hat{R} - \rho, F + \phi] \quad [12]$$

Adding a bias term to each value provides a direct way to break ties, and set default behaviors. Additional to our standard assumptions, if we also assume the fear/aversive term should take precedence over all other options when there is a tie, then $\phi > \rho > \eta$.

What about other values?. Another more mundane example would be to add a policy for a grooming drive, C .

$$\Pi^\pi = [\pi_E, \pi_R, \pi_C] \quad [13]$$

$$\operatorname{argmax}_{\Pi^\pi} [E - \eta, \hat{R} - \rho, C + \chi] \quad [14]$$

Where we assume the grooming term should take precedence over all other options, when there is a tie. That is, $\chi > \rho > \eta$.

What about fitting deterministic models?. Exploration in the lab is often described as probabilistic (?? ? ?). By definition probabilistic models do not make exact predictions of behavior, only statistical ones. Our approach is deterministic and so can make exact predictions. If, that is, we can fit it and it turns out to be correct. In species ranging from human to *Caenorhabditis elegans*, there are hundreds of exploration-exploitation experiments. A deterministic theory can, in principle, open up entirely new avenues for reanalysis. Testing and exploring these is a critical next step.

Does regret matter?. Another way around the dilemma is to ignore regret altogether. This is the domain of pure exploration methods, like the Gitten's Index (?), or probably-approximately correct approaches (?), or other bounded forms of reinforcement learning (?). Such methods can guarantee that you find the best reward value actions, eventually. These methods do not consider regret, as we have. But we argue that regret is critical to consider when resources and time are scarce, as they are in most natural settings.

Is the algorithmic run time practical?. It is common in computer science to study the run time of an algorithm as a measure of its efficiency. So what is the run time of our win stay, loose switch solution? There is unfortunately no fixed answer for the average or best case. The worst case algorithmic run time is linear and additive in the independent policies. If it takes T_E steps for π_E to converge, and T_R steps for π_R , then the worst case run time for π_π is $T_E + T_R$. Of course this is the worst case run time and would still be within reasonable ranges for learning in most naturalistic contexts.

What predictions does this theory make?.

Summary. We have argued one can put aside any intuitions that suggest a curious search is too inefficient, or open-ended, to be practical for reward collection. We have shown it can be made very practical, and mathematically optimal. We have shown there is a zero-regret way to solve the dilemma, a problem that has seemed unsolvable in those terms. Simply, be curious.

Acknowledgments

We wish to thank Jack Burgess, Matt Clapp, Kyle “Donovank” Dunovan, Richard Gao, Roberta Klatzky, Jayanth Koushik, Alp Muyesser, Jonathan Rubin, and Rachel Storer for their comments on earlier drafts. We also wishes to thank Richard Grant for his illustration work in Figure 1, and Jack Burgess for his assistance with Figure ??.

The research was sponsored by the Air Force Research Laboratory (AFRL/AFOSR) award FA9550-18-1-0251. The views and conclusions contained in this document are those of the authors and should not be interpreted as representing the official policies, either expressed or implied, of the Army Research Laboratory or the U.S. government. TV was supported by the Pennsylvania Department of Health Formula Award SAP4100062201, and National Science Foundation CAREER Award 1351748.

Mathematical Appendix.

Optimal substructure in memory. To find an optimal value solution for \hat{E} using the Bellman equation we must prove our memory \mathbf{M} has optimal substructure. This is because the normal route, which assumes the problem rests in a Markov Space, is closed to us. By optimal substructure we mean that the process of learning in \mathbf{M} , and therefore maximization of \hat{E} can be partitioned into a collection, or series of memories, each of which is itself an max value solution.

This opaque term of optimal substructure can be intuitively understood by looking in turn at another theoretical construct, Markov spaces.

In Markov spaces there are a set of states (S_0, S_1, \dots) , where the transition to the next S_t depends only on the previous state S_{t-1} . This limit means that if we were to optimize over these states, as in reinforcement learning, we know that we can treat each transition as its own “subproblem”, and ignore the history’s (S_0, S_1, \dots) effect on value, and therefore the overall situation has “optimal substructure”.

The problem for our definition of information value is it relies on memory which is necessarily composed of many past observations, in many orders, and so it cannot be a Markov space. So if we wish to use the Bellman equation, we need to find another way to establish optimal substructure. This is the focus of Theorem . In this theorem we implicitly assume that $(X = \mathbf{S}, \mathbf{A}, \mathbf{M}, f, \text{ and } \Lambda)$ are given, and that Λ is deterministic,

[Optimal substructure] If $V_{\pi_{\hat{E}}}^*$ is the optimal information value given by policy $\pi_{\hat{E}}$, a memory \mathbf{M}_t has optimal substructure if the last observation X can be removed from \mathbf{M} , by $\mathbf{M} - \mathbf{1}_t = f^{-1}(\mathbf{X}, \mathbf{M}_t)$ such that the resulting value $V_{t-1}^* = V_t^* - E_t$ is also optimal.

Proof. Given a known optimal value V^* given by $\pi_{\hat{E}}$ we assume for the sake of contradiction there also exists an alternative policy $\hat{\pi}_{\hat{E}} \neq \pi_{\hat{E}}$ that gives a memory $\hat{\mathbf{M}}_{t-1} \neq \mathbf{M}_{t-1}$ and for which $\hat{V}_{t-1}^* > V_{t-1}^*$.

To recover the known optimal memory \mathbf{M}_t we lift $\hat{\mathbf{M}}_{t-1}$ to $\mathbf{M}_t = f(\hat{\mathbf{M}}_{t-1}, \mathbf{X}_t)$. This implies $\hat{V}^* > V^*$ which in turn contradicts the purported original optimality of V^* and therefore $\hat{\pi}_{\hat{E}}$. \square

This proof requires two things. First, we need to a mechanism of forgetting, of a very particular kind. We assume that the last learning step $f(\mathbf{X}, \mathbf{M}) \rightarrow \mathbf{M}'$ can be undone by a new vector valued function f^{-1} , such that $f(\mathbf{X}, \mathbf{M}') \rightarrow \mathbf{M}$. In other words we must assume what was last remembered, can always be exactly forgotten. Second, we assume the environmental dynamics, given by the transition function Λ , are deterministic.

Determinism is consistent with the natural world which does evolve in a deterministic way, at the scale of animal behavior that we concerned with at least. This assumption is however at odds with much of reinforcement learning theory (?) and past experimental work. For example, (?). Both tend to study stochastic environments. We addressed this discrepancy in the main text, using numerical simulations.

Bellman optimal information collection. We aim to find a series of actions $(\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_T)$, drawn from a set \mathbb{A} , that maximize each payout $(\hat{E}_0, \hat{E}_1, \dots, \hat{E}_T)$ so the total payout received is as large as possible. If there is a policy $\mathbf{A} = \pi(\mathbf{X})$ to take actions, based on a series of observations (X_0, X_1, \dots, X_T) , given

by Λ , then an optimal policy π^* will always find the maximum total value $V^* = \operatorname{argmax}_{\mathbf{A} \in \mathbb{A}} \sum_T \hat{E}_t$. In the form above one would need to reconsider the entire sequence of actions for any one change to that sequence, leading to a combinatorial explosion. Bellman’s insight was a way to make this last problem simpler by breaking it down into a small set of subproblems that we can solve in a tractable way without an explosion in complexity. This is his principle of optimality, which reads:

An optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with regard to the state resulting from the first decision. (?)

Mathematically Bellman’s principle allows us to translate the full problem, $V^* = \sum_T r_1, r_2, \dots, r_T$ to a recursive one. Having already proven the optimal substructure of \mathbf{M} , and given an arbitrary starting value E_0 , the Bellman solution to curiosity optimization of \hat{E} is therefore given by,

$$V_{\hat{E}}^*(\mathbf{S}) = \operatorname{argmax}_{\mathbf{A} \in \mathbb{A}} \left[\hat{E}_t + V_{\hat{E}}^*(\Lambda(\mathbf{S}, \mathbf{A})) \right] \quad [15]$$

Optimal exploration. Recall from the main text we consider that a good exploration should,

1. Exploration should visit all available states of the environment at least once.
2. Exploration should cease when learning has plateaued.
3. Exploration should take as few steps as possible to achieve 1 and 2.

If $\pi_{\hat{E}}$ to a deterministic policy this makes proving these three properties amounts to solving sorting problems on \hat{E} , and assuming $\mathbf{X} = \mathbf{S}$. That is, if a state is visited by our algorithm it must have the highest value, by definition. So if every state must be visited every state must, at one time or another, give the maximum value. This is certain to happen if we know that all values will begin contracting towards zero, in finite time.

Definitions. Let \mathbb{Z} be the set of all visited states, where \mathbb{Z}_0 is the empty set $\{\}$ and \mathbb{Z} is built iteratively over a path P , such that $\mathbb{Z} \rightarrow \{\mathbf{S} | \mathbf{S} \in \mathbb{S} \text{ and } x \notin \mathbb{Z}\}$.

To formalize the idea of ranking we take an algebraic approach. Give any three real numbers (a, b, c) ,

$$a \leq b \Leftrightarrow \exists c; b = a + c \quad [16]$$

$$a > b \Leftrightarrow (a \neq b) \wedge (b \leq a) \quad [17]$$

Theorem 1 (Complete exploration). *Given some arbitrary value \hat{E}_0 , an exploration policy governed by $\pi_{\hat{E}}$ will visit all states $\mathbf{S} \in \mathbb{S}$ in finite number of steps T .*

Proof. Let $\mathbf{E} = (\hat{E}_1, \hat{E}_2, \dots)$ be ranked series of \hat{E} values for all states \mathbf{S} , such that $(\hat{E}_1 \geq \hat{E}_2 \geq \dots)$. To swap any pair of values $(\hat{E}_i \geq \hat{E}_j)$ so $(\hat{E}_i \leq \hat{E}_j)$ by Eq. 16 $\hat{E}_i - c = \hat{E}_j$.

Therefore, again by Eq. 16, $\exists \delta \hat{E} \rightarrow -c$.

Recall: Axiom 5 - $\nabla^2 \mathbf{M} < 0$ after a finite time T^* .

However if we wished to instead swap $(\hat{E}_i \leq \hat{E}_j)$ so $(\hat{E}_i \geq \hat{E}_j)$ by definition $\nexists c; \hat{E}_i + c = \hat{E}_j$, as $\nexists \delta \hat{E} \rightarrow c$.

To complete the proof, assume that some policy $\hat{\pi}_{\hat{E}} \neq \pi_E^*$. By definition policy $\hat{\pi}_{\hat{E}}$ can be any action but the maximum, leaving $k - 1$ options. Eventually as $t \rightarrow T^*$ the only possible swap is between the max option and the k th, but as we have already proven this is impossible as long as Axiom 5 holds. Therefore, the policy $\hat{\pi}_{\hat{E}}$ will leave at least 1 option unexplored and $S \neq Z$. \square

Theorem 2 ([Efficient exploration]). *An exploration policy governed by $\pi_{\hat{E}}^*$ will revisit all states in exact proportion to their information value.*

Proof. Recall: Theorem 1. Recall: π_E^ is a maximum value deterministic algorithm. Recall: Axiom 2. Each time π_E^* visits a state S , so $M \rightarrow M'$, and after T^* it is axiomatically true $\hat{E}' < \hat{E}$*

By induction then, if π^*E will visit all $S \in \mathbb{S}$ in T^* trials, it will revisit them at most $2T^*$, therefore as $t \rightarrow \infty$, $E \rightarrow \eta$, where $\eta > 0$. \square

These exploration proofs come with some fine print, for practical work. E_0 can be any positive and finite real number, $E_0 > 0$. Different choices for E_0 will not change the proofs, especially their convergence. So in that sense one can choose it in an arbitrary way. Different choices for E_0 can however change individual choices, and their order. This can be quite important in practice, especially when trying to describe some real data.

Optimality of Π_π . Recall that in the main text we introduce the equation below as a candidate with zero regret solution to exploration-exploitation problems, set in terms the mixed value sequence $V_{\hat{E}R}$.

$$\Pi_\pi = \begin{cases} \pi_{\hat{E}}^* & : \hat{E} - \eta > R + \rho \\ \pi_R & : \hat{E} - \eta < R + \rho \end{cases} \quad [18]$$

In the following section we prove two things about the optimality of π_π . First, if π_R had any optimal asymptotic property for value learning before their inclusion into our scheduler, they retain that optimal property under π_π when $\eta = 0$, or is otherwise sufficiently small. Second, show that if both π_R and π_E are greedy, and π_π is greedy in its definition, then Eq 9 is certain to maximize total value. The total value of R and \hat{E} is the exact quantity to maximize if information seeking and reward seeking are equally important, overall. This is, as the reader may recall, one of our key assumptions. Proving this optimality is analogous to the classic activity selection problem from the job scheduling literature (??).

Theorem 3 (π_π is unbiased). *Let any S, A, M, π_R, π_E , and δ be given. Assuming an infinite time horizon, if π_E is optimal and π_R is optimal, then π_π is also optimal in the same sense as π_E and π_R .*

Proof. The optimality of π_π can be seen by direct inspection. If $p(R = 0) > 0$ we are given an infinite horizon, then π_E will have a unbounded number of trials meaning the optimality of P^* holds. Likewise, $\sum E < \eta$ as $T \rightarrow \infty$, ensuring π_R will dominate π_π therefore π_R will asymptotically converge to optimal behavior. \square

In proving this optimality of π_π we limit the probability of a positive reward to less than one, denoted by $p(R_t = 1) < 1$. Without this constraint the reward policy π_R would always dominate π_π when rewards are certain. While this might be useful in some circumstances, from the point of view π_E it is extremely suboptimal. The model would never explore. Limiting $p(R_t = 1) < 1$ is a reasonable constraint, as rewards in the real world are rarely certain. A more naturalistic way to handle this edge case is to introduce reward satiety, or a model physiological homeostasis (??).

In classic scheduling problems the value of any job is known ahead of time (??). In our setting, this is not true. Reward value is generated by the environment, *after* taking an action. In a similar vein, information value can only be calculated *after* observing a new state. Yet Eq. 9 must make decisions *before* taking an action. If we had a perfect model of the environment, then we could predict these future values accurately with model-based control. In the general case though we don't know what environment to expect, let alone having a perfect model of it. As a result, we make a worst-case assumption: the environment can arbitrarily change-bifurcate-at any time. This is a highly nonlinear dynamical system (??). In such systems, myopic control-using only the most recent value to predict the next value-is known to be an robust and efficient form of control (??). We therefore assume that last value is the best predictor of the next value, and use this assumption along with Theorem 3 to complete a trivial proof that Eq. 9 maximizes total value.

A win-stay, lose-switch solution. If we prove π_π has optimal substructure, then using the same replacement argument (??) as in Theorem 3, a greedy policy for π_π will maximize total value.

Theorem 4 (No regret - mixed values). *When either $\pi_{\hat{E}}$ or π_R is in control under Π_π , all actions are zero regret in terms of $V_{\hat{E}R}$. That is, $\sum_{k=0}^T G = 0$.*